

# **Interfacing MicroC/OS to TCP/IP**

# Interfacing Small Systems to TCP/IP

---

- TCP/IP is a complicated protocol that requires substantial memory space and program complexity.
- Special care is required when interfacing small systems, with limited memory and CPU resources, to TCP/IP.
- As implementation examples, we will consider how MicroC/OS has been interfaced to two TCP/IP stacks:
  1. NetBurner's TCP/IP stack (the ECE 315 lab system)
  2. The Lightweight IP (lwIP) portable, open-source TCP/IP stack. [Adam Dunkels, "Design and Implementation of the LWIP TCP/IP Stack," Swedish Institute of Computer Science, Feb. 20, 2001.]

# Opportunities for Saving Time and Space

---

- Strict partitioning of the TCP/IP interface according to the protocol layers can be inefficient.  
*Ex.* It may be more efficient at times for TCP to be able to look directly into the IP header without always having to go through the IP layer software.
- We want to avoid unnecessary copying of data and packet buffers between the application and the communication subsystem, or between different layers in the protocol stack.  
*Ex.* We may want the application to use the same buffers as the communication subsystem even if this violates strict partitioning between applications and the TCP/IP stack and operating system.

# Implement the Protocol Layers as Tasks?

---

- One could employ separate tasks for each protocol layer:
  - (4) Application task(s)
  - (3) Transport (TCP) layer task
  - (2) Internetworking (IP) layer task
  - (1) Device driver task (for Ethernet, RS-232C, etc.)

This approach was taken by NetBurner's TCP/IP stack.
- *Advantages:*
  - interlayer interfaces strictly defined and enforceable
  - debugging and code maintenance is easier
  - can more easily modify protocol settings at run-time
- *Disadvantages:*
  - larger required number of task context switches
  - need to carefully control access to shared buffers
  - access to multi-layer protocol data is slowed down

# Some Other Implementation Options

---

- Embed the TCP/IP stack inside the kernel. User tasks must access the stack using system calls (also called “traps”).

*Advantages:* Fast & efficient; users shielded from stack details.

*Disadvantages:* More difficult to change or modify the stack.

- Combine protocol layers into one task, outside the kernel. User code may, in some cases, be combined with this same task.

*Advantages:* Protocol stack is more portable across kernels.

Opportunities to gain efficiency in space & time.

*Disadvantages:* Stack software is less clearly partitioned apart from user code.

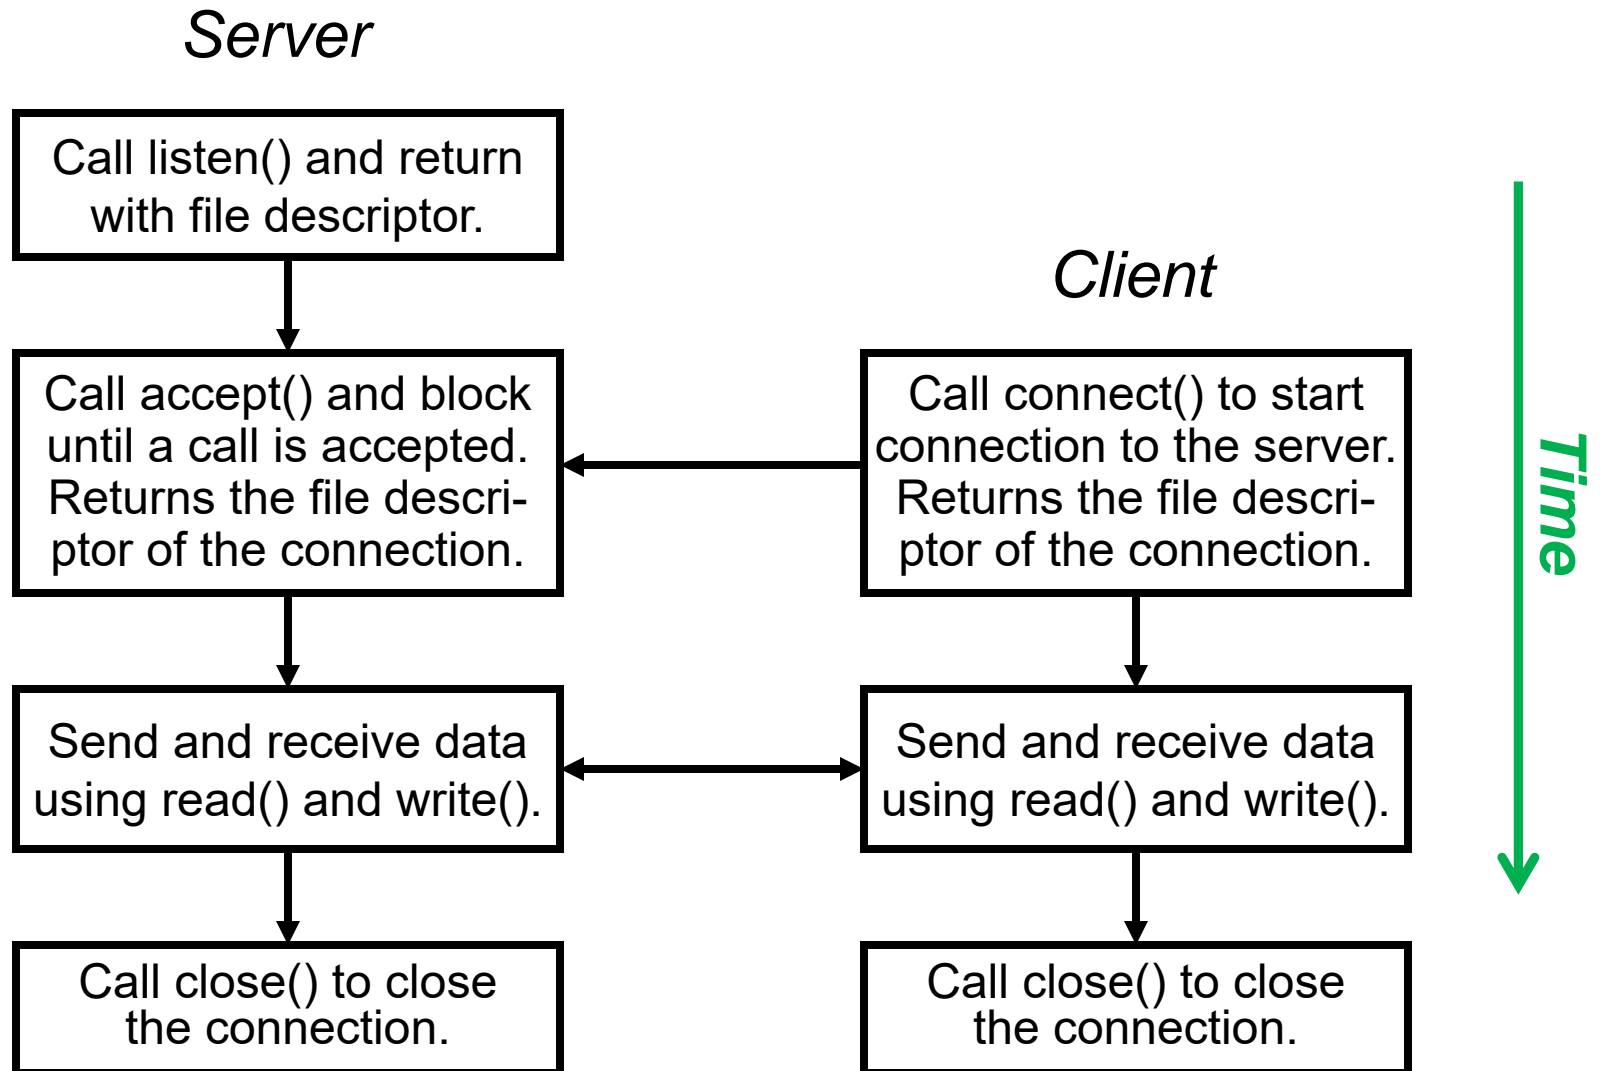
- Lightweight IP takes the last approach, with the option of keeping user code in the same or a separate task from the stack task.

# The NetBurner TCP/IP Stack in MicroC/OS

---

- NetBurner's TCP/IP stack is partitioned across several tasks that handle the different protocol layers:
  - Priority 38: Ethernet hardware (FEC) driver task
  - Priority 39: IP entity task
  - Priority 40: TCP entity task
  - Priority 44: PPP entity task (for TCP/IP over RS-232C)
  - Priority 45: HTTP web server task (created by StartHTTP)
  - Priority 50: Default priority of the UserMain() task
- The FEC driver task has the highest priority (i.e., lowest priority number in MicroC/OS) so that Ethernet hardware actions will be handled promptly before IP and TCP.

# NetBurner's Client-Server Architecture



# TCP/IP Application Programming Interface (1)

---

New data types introduced in the TCP/IP stack:

- `IPADDR` - data structure that holds one IP address or mask

File descriptor error return codes (all negative values):

- `TCP_ERR_TIMEOUT` - connection timed out
- `TCP_ERR_NOCON` - connection not yet established
- `TCP_ERR_CLOSING` - connection closed and not available
- `TCP_ERR_NOSUCH_SOCKET` - socket does not exist
- `TCP_ERR_NONE_AVAIL` - no more sockets are available
- `TCP_ERR_CON_RESET` - connection reset by other side
- `TCP_ERR_CON_ABORT` - internal stack error



# TCP/IP Application Programming Interface (2)

---

```
void InitializeStack( IPADDR IP_addr = 0,  
                    IPADDR IP_mask = 0,  
                    IPADDR IP_gateway = 0 );
```

- function that initializes the TCP/IP stack
- typically called by UserMain()
- "IP\_addr" is the IP address of the local node
- "IP\_mask" is the IP mask of the local node
- "IP\_gateway" is the IP address of the gateway for the local node
- If called with no inputs, then default IP values are copied from the system configuration record.

```
void KillStack( );
```

- function that shuts down the TCP/IP stack
- typically called by UserMain()

# TCP/IP Application Programming Interface (3)

---

**int listen( IPADDR IP\_addr, WORD port\_num, BYTE max\_pends );**

- called by the server to put the TCP entity into the LISTEN state
- “IP\_addr” specifies a client IP; value INADDR\_ANY means that a connection/call will be accepted from any client IP
- “port\_num” specifies the listening port number on the server
- “max\_pends” specifies maximum number of pending connections
- return value is the listening file descriptor, i.e. a dummy socket

**int accept( int fdl, IPADDR \*Pclient\_IP, WORD \*Pclient\_port,  
WORD timeout );**

- called by the server to block until a client call is received
- “fdl” is the listening file descriptor, i.e. the dummy listening socket
- “Pclient\_IP” points to the IP address of a new calling client
- “Pclient\_port” points to the port number on the client
- Pclient\_IP and Pclient\_port can be left null if information not wanted
- “timeout” is the maximum number of ticks waiting for a call;  
a “timeout” value of 0 means wait forever for the next call
- return value is the network file descriptor, i.e. the working socket

# TCP/IP Application Programming Interface (4)

---

**int read( int fdnet, char \*buf, int buf\_siz );**

- used by servers & clients to receive a char array payload
- “fdnet” is the network file descriptor, i.e. the working socket
- “buf” stores the retrieved payload, which is a char array
- “buf\_siz” is the maximum number of chars to be read
- positive return value is the actual number of new chars in "buf"
- negative return code means the socket was closed by the other side.

**int write( int fdnet, char \*buf, buf\_siz );**

- used by servers & clients to send a char array payload
- “fdnet” is the network file descriptor, i.e. the working socket to write
- “buf” holds the payload to be written, which is a char array
- “buf\_siz” is the number of chars to be written
- return code is the number of chars actually written

# TCP/IP Application Programming Interface (5)

---

**int ReadWithTimeout( int fdnet, char \*buf, int buf\_siz,  
                          unsigned long timeout );**

- used by servers & clients to receive a char array payload
- “fdnet” is the network file descriptor, i.e. the working socket
- “buf” stores the received payload, which is a char array
- “buf\_siz” is the maximum number of chars to be read
- “timeout” is a timeout to socket closure, expressed in timer ticks
- a positive return value is the number of chars in char\_buf
- a negative return value means the other party closed the socket

**void writestring( int fdnet, char \*buf );**

- an alternative to the write() function
- “fdnet” is the network file descriptor, i.e. the working socket
- writes a null-terminated string to a network file descriptor

# TCP/IP Application Programming Interface (6)

---

**int connect( IPADDR remote\_IP, WORD local\_port, WORD remote\_port, DWORD timeout );**

- connects a client to the specified port on a remote server
- “remote\_IP” specifies the IP address of the remote server
- “local\_port” specifies the port number to use for the connection; a value of 0 causes the stack to choose an unused port
- “remote\_port” specifies the port number on the server
- “timeout” is a timeout, expressed in timer ticks; a value of 0 means that the function will wait forever
- return value is the network file descriptor, i.e. the working socket

**void close( fdnet );**

- called by both servers & clients to end a TCP connection
- “fdnet” is the network file descriptor, i.e. the working socket

# Example: A Simple Echo Server (1)

---

```
#include "predef.h"
#include <stdio.h>
#include <ctype.h>
#include <startnet.h>
#include <autoupdate.h>
#include <dhcpcclient.h>
#include <ip.h>
#include <tcp.h>

#define ECHO_LISTEN_PORT 23    // use the TELNET port
#define RX_BUFSIZE 4096

DWORD EchoServerTaskStack[USER_TASK_STK_SIZE];
                        __attribute__((aligned(4)));

const char *Echo_server = "Simple echo server";

char RXBuffer[RX_BUFSIZE];
```

# Example: A Simple Echo Server (2)

---

```
void EchoServerTask( void *pd )
{
    IPADDR client_IP;
    WORD port;
    int ListenPort = *(int *)pd;
    int fdListen = listen(INADDR_ANY, ListenPort, 7);
    if (fdListen > 0) {
        while (1) {
            int fdnet = accept(fdListen, &client_IP, &port, 0);
            while (fdnet > 0) {
                writestring( fdnet, "Welcome to the server\r\n" );
                int n = 0;
                do {
                    n = read( fdnet, RXBuffer, RX_BUFSIZE );
                    write( fdnet, RXBuffer, n );
                } while ((n > 0) && (RXBuffer[0] != '.'))
                close( fdnet );
                fdnet = 0;
            } // end while (fdnet > 0)
        } // end while (1)
    } // if (fdListen > 0)
}
```

# Example: A Simple Echo Server (3)

---

```
void UserMain( void *pd)
{
    InitializeStack();

    if (EthernetIP == 0) { GetDHCPAddress(); }
    EnableAutoUpdate();

    OSChangePrio(MAIN_PRIO);

    OSTaskCreate( EchoServerTask,
        (void *) ECHO_LISTEN_PORT,
        (void *) &EchoServerTaskStack[USER_TASK_STK_SIZE],
        (void *) EchoServerTaskStack,
        MAINPRIO-1) ;    // higher priority than UserMain()

    while (1) {
        OSTimeDly( TICKS_PER_SECOND * 5 );
    } // end while (1)
}
```



# NetBurner's HTTP Server (1)

---

- The simple echo server could be used to build an HTTP server; however, NetBurner already provides an extensible HTTP server.
- NetBurner's HTTP server (a “webserver”) is implemented as a task at priority HTTP\_PRIO, which is 45 by default.
- Required steps to use NetBurner's webserver:
  - include the header file “http.h”
  - include the header file “htmlfiles.h”
  - initialize the TCP/IP stack by calling “InitializeStack()”
  - start up the HTTP server by calling “StartHTTP()”
- A default web page will be returned for GET requests to port 80.
- The HTTP server is shut down by calling “StopHTTP()”.

# NetBurner's HTTP Server (2)

---

New types that are defined and used in the HTTP server:

```
typedef char *      PSTR    // pointer to a string
typedef const char * PCSTR  // pointer to a constant string
```

**void StartHTTP( WORD port = 80 );**

- starts up the HTTP server task
- default HTTP port of 80 can be overridden
- usually called by UserMain();
- must have previously called InitializeStack();

**void StopHTTP();**

- shuts down the HTTP server task

# NetBurner's HTTP Server (3)

---

```
// Default simple handler of HTTP GET requests
int BaseDoGet( int sock, PSTR url, PSTR rxBuffer )
{
    if ( *url == 0 ) {    // null url, so GET default webpage
        RedirectResponse( sock, url_of_default_webpage );
        // url_of_default_webpage is typically "HTML/index.htm"
        // All other webserver .htm files are in the HTML dir.
        return 1;
    }
    if ( httpstricmp(url,"ECHO") ) {    // ECHO server loop
        while ( *url==0 ) url++;        // Advance past nul's
        SendTextHeader( sock );        // Send HTML header
        *url = ' ';                    // Insert a space char
        writestring( sock, rxBuffer ); // Echo back the message
        return 1;
    }
    if ( !SendFullResponse( url, sock ) ) { // GET webpage
        NotFoundResponse( sock, url );    // else 404 not found
    }
    return 0;
}
```

# NetBurner's HTTP Server (4)

---

**void RedirectResponse( int fd, PCSTR url );**

- redirect the current GET request to the webpage at "url"
- send the webpage to socket "fd"

**int SendFullResponse( PCSTR url, int fd );**

- responding by GET-ing the stored webpage for the given "url"
- the webpage must have been stored previously in the html subdirectory of the project directory before building the project
- the webpage may have dynamic parts that change at runtime
- send the HTTP header plus webpage back to socket "fd"
- returns 1 if webpage was found; otherwise, returns 0

**void NotFoundResponse( int fd, PCSTR url );**

- no webpage was found on the server with the given "url"
- send back an HTTP "not found" response to socket "fd"

# NetBurner's HTTP Server (5)

---

**int httpstricmp( PCSTR string, PCSTR pattern );**

- "string" contains an array of chars
- "pattern" contains an array of UPPER CASE chars
- returns 1 if the shorter "string" or "pattern" appears as a prefix of the longer array of "string" or "pattern"; otherwise, returns 0

**void SendTextHeader( int fd );**

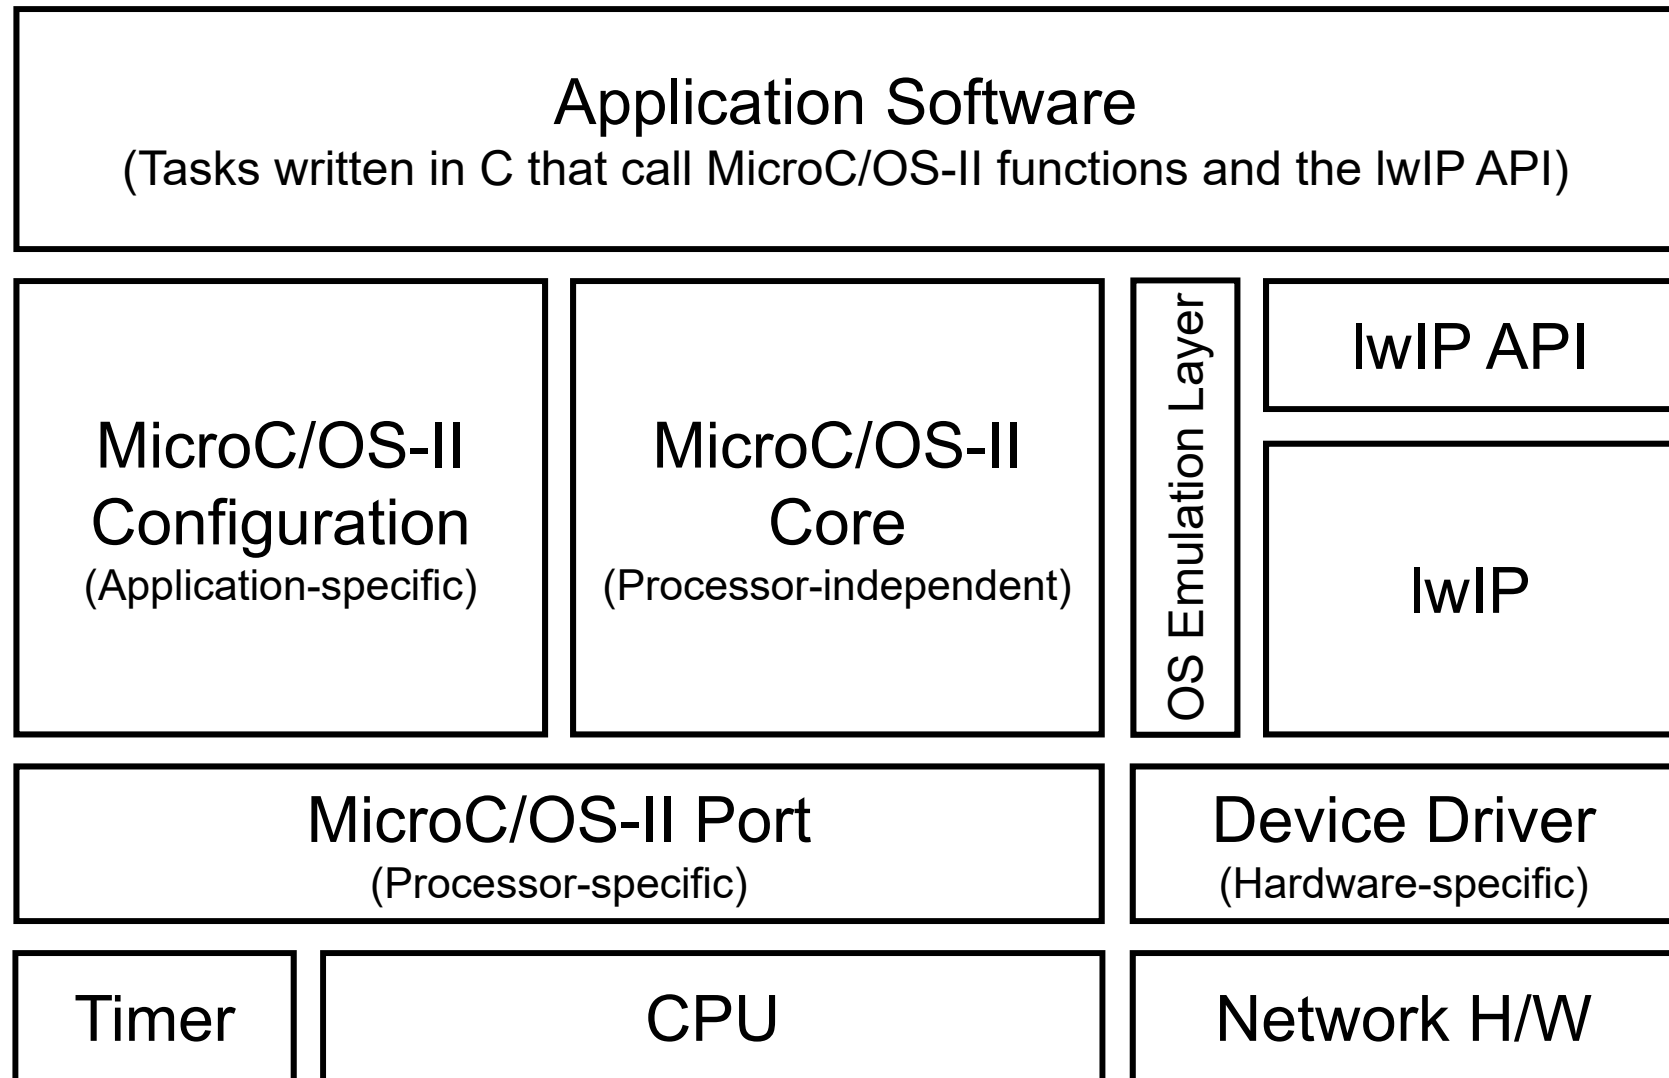
- send the appropriate HTTP response header for a text reply
- must be the first part of the HTTP response to socket "fd"
- followed by a write of a plain text file
- used to response to HTTP ECHO requests

**void SendHTMLHeader( int fd );**

- send the appropriate HTTP response header for an HTML file
- followed by a write of an HTML file

# Architecture showing MicroC/OS-II with lwIP

---



# Components of Lightweight IP

---

## lwIP Application Program Interface (API):

- simplified function calls for accessing TCP/IP network
- very similar to the socket API in Unix

## lwIP Process:

- kernel/OS-independent implementation of TCP/IP

## Operating System Emulation Layer:

- minimal set of functions required of the O.S. by lwIP

## Network Hardware Interface:

- provides access to physical layer: Ethernet, RS-232C, Bluetooth, wireless LAN (IEEE 802.11b, 802.11a), etc.

## Device Driver:

- software that initializes and operates the hardware

# Operating System Emulation Layer

---

- Operating System specific functions are kept together in this layer to make it easier to port lwIP to different OS's.
- Required functionality in the operating system or kernel:
  - One-shot timer with at least 200 ms resolution
  - Semaphore process synchronization mechanism
  - Mailbox system with:
    - (1) nonblocking “post to message” queue function
    - (2) blocking “receive from message” queue function
- Lightweight IP has its own buffer management system so that it can exploit potential efficiencies and also be more independent of the other parts of the system.

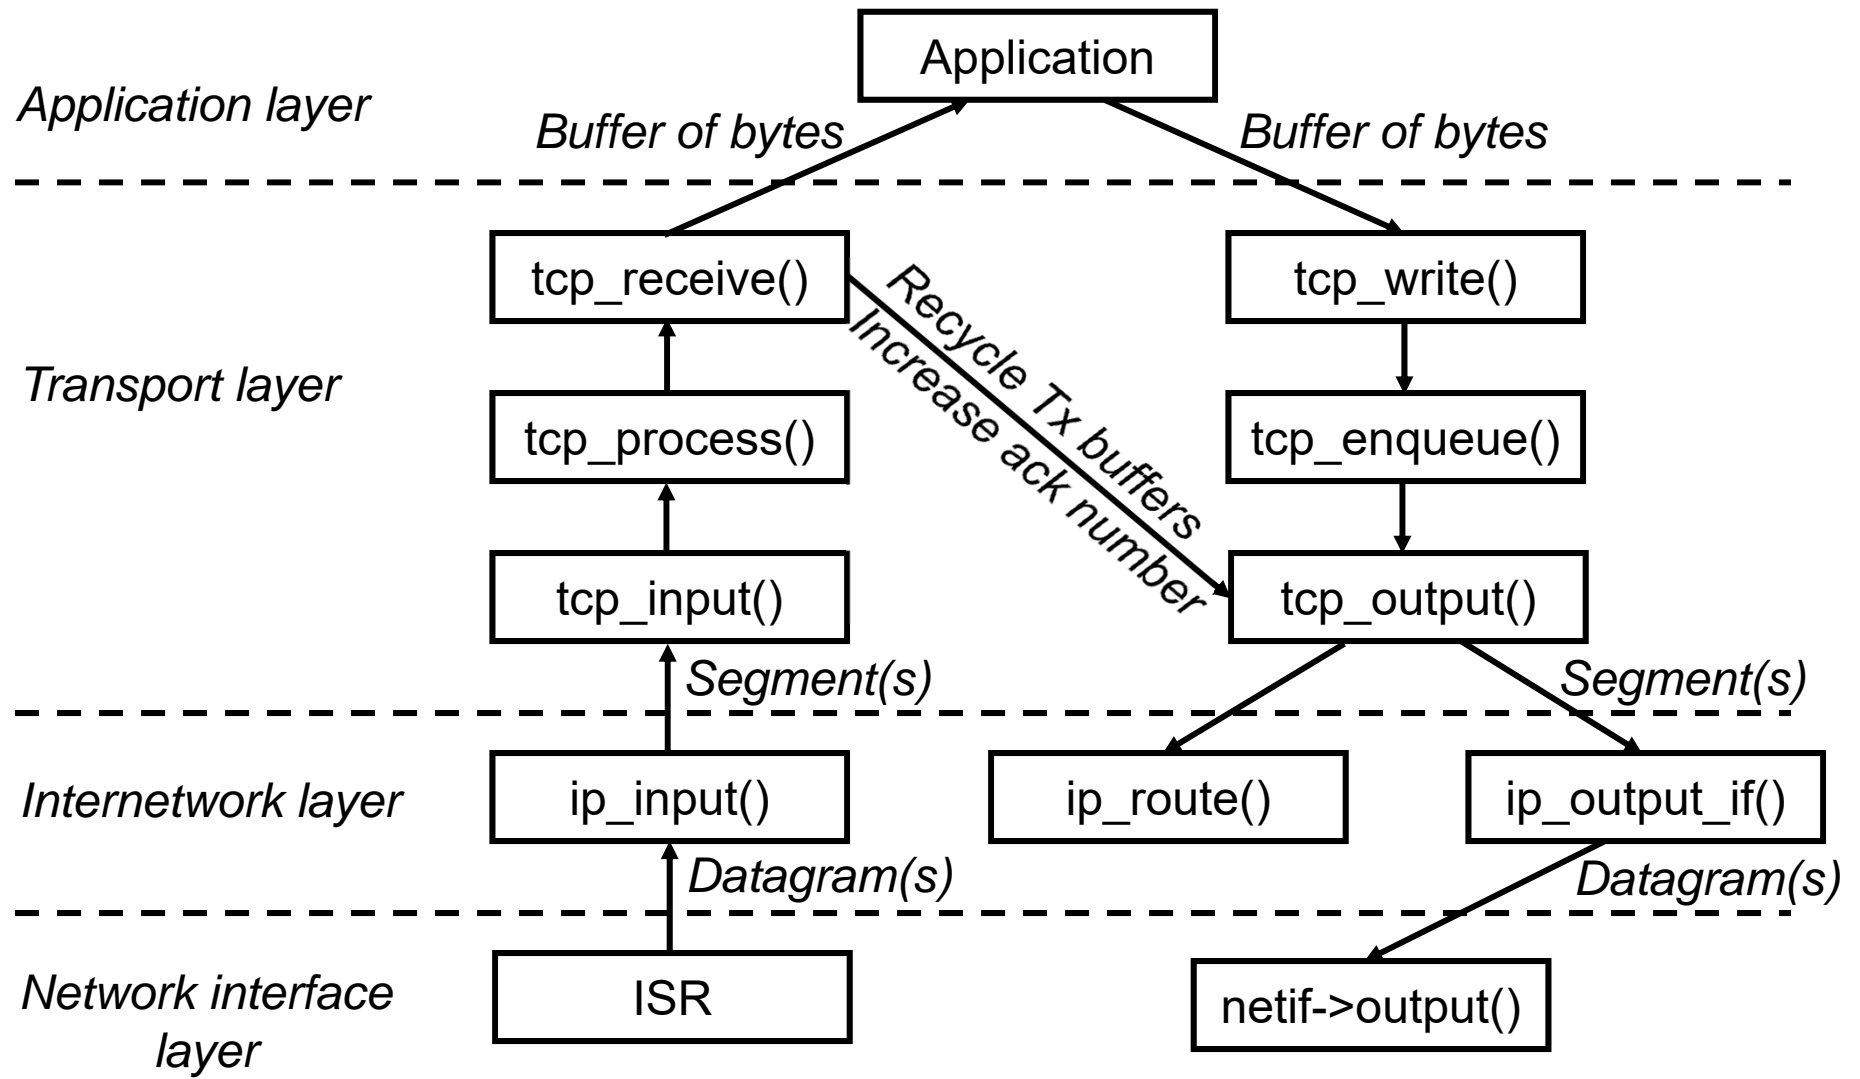


# Device Driver

---

- A ***device driver*** is a body of software that takes care of initializing and operating a hardware device or interface. Typical components in a device driver:
  - initialization routine
  - data output routine
  - data input routine
  - interrupt service routine(s)
  - data buffers, state variables
  - query interface status
- In lwIP, as in Unix systems, there can be several device drivers available for operating different network interfaces.
- All device drivers can be used in the same way by higher-level lwIP software, so lwIP is shielded from device-specific details.

# TCP/IP Data Processing in lwIP



# Processing of Incoming Datagrams in lwIP

---

- `ip_input()`
- verifies the IP header of the datagram
  - datagram discarded if error detected
  - verifies if datagram destination is host node
  - passes resulting segment up to the appropriate higher layer protocol (TCP, UDP, ICMP)
- `tcp_input()`
- verifies TCP header checksum in segments
  - parses options in the TCP header
  - from port number, determines TCP connection
- `tcp_process()`
- make any state transitions in TCP protocol
  - call `tcp_receive()` if connection ready for data
- `tcp_receive()`
- strip off TCP header and pass data to application
  - recycle any ACKed outgoing segments in buffer
  - call `tcp_output()` if ACKed transmitted bytes allow more outgoing segments to be transmitted

# Processing of Outgoing Data Bytes in lwIP

---

- `tcp_write()` - called to pass data byte buffer to TCP layer
- `tcp_enqueue()` - if necessary, break up data buffer into pieces
  - queue up the resulting TCP segments
- `tcp_output()` - check to see if more segments can be sent
  - if yes, call `ip_route` then `ip_output_if()`
- `ip_route()` - find the next node in the datagram's route and select the outgoing port from the present node
- `ip_output_if()` - assemble datagram, and call device driver
- `ifnet->output()` - the device driver that loads the datagram into the transmitter hardware